



UNIVERSITÀ  
DEGLI STUDI  
DI PADOVA




DIPARTIMENTO  
**MATEMATICA**

DIPARTIMENTO DI MATEMATICA "TULLIO LEVI-CIVITA"

STATISTICAL LEARNING FINAL PROJECT

# Employee Attrition Classification



AUTHORS

**Zeynep TUTAR - 2106038**

**Aysenur Oya ÖZEN - 0000000**

SUPERVISOR

**Prof. Alberto ROVERATO**

**Academic Year:  
2023/2024**

## Contents

<b>Introduction to Dataset</b>	<b>2</b>
Description of the Features . . . . .	2
<b>Data Analysis</b>	<b>3</b>
Data Preprocessing . . . . .	5
Categorical Features . . . . .	6
Target Values . . . . .	10
Numeric Features . . . . .	11
Outliers . . . . .	13
Features vs. Target . . . . .	16
Categorical Features vs. Target . . . . .	16
Numerical Features vs. Target . . . . .	16
Correlation Matrix . . . . .	16
Partial Correlation Matrices . . . . .	18
<b>Data Preparation</b>	<b>18</b>
Handling Categorical Features . . . . .	18
Normalization . . . . .	19
Train-Test-Split . . . . .	19
<b>Predictive Classification Models</b>	<b>20</b>
Logistic Regression . . . . .	21
Basic Logistic Classifier . . . . .	21
Logistic Regression with Backward Stepwise Search . . . . .	23
Logistic Regression with Shrinkage Method . . . . .	25
Comparison of Logistic Classifiers . . . . .	26
Another Classification Model . . . . .	29
<b>Model Results</b>	<b>29</b>
Performance Metrics and Confusion Matrix . . . . .	29

## Introduction to Dataset

The aim of this project is to develop two predictive models to determine employee attrition of a company. The dataset<sup>1</sup> used for this project is a simulated dataset designed for the analysis and prediction of employee attrition. It contains detailed information about various aspects of an employee's profile, including demographics, job-related features, and personal circumstances. The dataset contains 74,498 samples. Each record includes a unique Employee ID and features that influence employee attrition. The goal is to understand the factors contributing to attrition and develop predictive models to identify at-risk employees.

The dataset is already split into train and test but in order to better understand the data, it is crucial to analyse the dataset as a whole.

```
# import the train and test datasets
data_train <- read.csv("data/train.csv", stringsAsFactors = TRUE)
data_test <- read.csv("data/test.csv", stringsAsFactors = TRUE)

# merge the datasets
data <- rbind(data_train, data_test)
attach(data)
```

## Description of the Features

The features of the dataset are presented below:

- **Employee ID:** A unique identifier assigned to each employee.
- **Age:** The age of the employee, ranging from 18 to 60 years.
- **Gender:** The gender of the employee
- **Years at Company:** The number of years the employee has been working at the company.
- **Monthly Income:** The monthly salary of the employee, in dollars.
- **Job Role:** The department or role the employee works in, encoded into categories such as Finance, Healthcare, Technology, Education, and Media.
- **Work-Life Balance:** The employee's perceived balance between work and personal life, (Poor, Below Average, Good, Excellent)
- **Job Satisfaction:** The employee's satisfaction with their job: (Very Low, Low, Medium, High)
- **Performance Rating:** The employee's performance rating: (Low, Below Average, Average, High)
- **Number of Promotions:** The total number of promotions the employee has received.
- **Distance from Home:** The distance between the employee's home and workplace, in miles.
- **Education Level:** The highest education level attained by the employee: (High School, Associate Degree, Bachelor's Degree, Master's Degree, PhD)
- **Marital Status:** The marital status of the employee: (Divorced, Married, Single)

<sup>1</sup><https://www.kaggle.com/datasets/stealthtechnologies/employee-attrition-dataset/data>

- **Job Level:** The job level of the employee: (Entry, Mid, Senior)
- **Company Size:** The size of the company the employee works for: (Small,Medium,Large)
- **Company Tenure:** The total number of years the employee has been working in the industry.
- **Remote Work:** Whether the employee works remotely: (Yes or No)
- **Leadership Opportunities:** Whether the employee has leadership opportunities: (Yes or No)
- **Innovation Opportunities:** Whether the employee has opportunities for innovation: (Yes or No)
- **Company Reputation:** The employee's perception of the company's reputation: (Very Poor, Poor,Good, Excellent)
- **Employee Recognition:** The level of recognition the employee receives:(Very Low, Low, Medium, High)
- **Attrition:** Whether the employee has left the company, encoded as 0 (stayed) and 1 (Left).

## Data Analysis

In order to develop predictive models, first it is necessary to perform exploratory data analysis (EDA) and modify the format of the data if necessary.

```
# installing required libraries
```

```
library(class)
library(e1071)
library(car)
library(corrplot)
library(glmnet)
library(dplyr)
library(pROC)
library(knitr)
library(leaps)
```

```
# Descriptive statistics of DataFrame
```

```
summary(data)
```

Employee.ID	Age	Gender	Years.at.Company
Min. : 1	Min. :18.00	Female:33672	Min. : 1.00
1st Qu.:18625	1st Qu.:28.00	Male :40826	1st Qu.: 7.00
Median :37250	Median :39.00		Median :13.00
Mean :37250	Mean :38.53		Mean :15.72
3rd Qu.:55874	3rd Qu.:49.00		3rd Qu.:23.00
Max. :74498	Max. :59.00		Max. :51.00
Job.Role	Monthly.Income	Work.Life.Balance	Job.Satisfaction
Education :15658	Min. : 1226	Excellent:13432	High :37245
Finance :10448	1st Qu.: 5652	Fair :22529	Low : 7457
Healthcare:17074	Median : 7348	Good :28158	Medium :14717
Media :11996	Mean : 7299	Poor :10379	Very High:15079

```

Technology:19322    3rd Qu.: 8876
                  Max.    :16149

    Performance.Rating Number.of.Promotions Overtime    Distance.from.Home
Average      :44719    Min.      :0.0000    No :50157    Min.      : 1.00
Below Average:11139    1st Qu.:0.0000    Yes:24341    1st Qu.:25.00
High        :14910    Median :1.0000                Median :50.00
Low         : 3730    Mean   :0.8329                Mean   :49.99
                  3rd Qu.:2.0000                3rd Qu.:75.00
                  Max.    :4.0000                Max.    :99.00

    Education.Level    Marital.Status    Number.of.Dependents    Job.Level
Associate Degree :18649    Divorced:11078    Min.      :0.00    Entry :29780
Bachelor's Degree:22331    Married :37419    1st Qu.:0.00    Mid   :29678
High School      :14680    Single  :26001    Median :1.00    Senior:15040
Master's Degree  :15021                Mean   :1.65
PhD              : 3817                3rd Qu.:3.00
                  Max.    :6.00

Company.Size    Company.Tenure    Remote.Work    Leadership.Opportunities
Large :14912    Min.      : 2.00    No :60300    No :70845
Medium:37231    1st Qu.: 36.00    Yes:14198    Yes: 3653
Small :22355    Median : 56.00
                  Mean   : 55.73
                  3rd Qu.: 76.00
                  Max.    :128.00

Innovation.Opportunities    Company.Reputation    Employee.Recognition
No :62394                    Excellent: 7414    High      :18550
Yes:12104                    Fair      :14786    Low       :29620
                               Good       :37182    Medium    :22657
                               Poor       :15116    Very High: 3671

```

Attrition  
Left :35370  
Stayed:39128

```
# Data types of columns
```

```
str(data)
```

```

'data.frame':   74498 obs. of  24 variables:
 $ Employee.ID      : int  8410 64756 30257 65791 65026 24368 64970 36999 32714 15944 ...
 $ Age              : int   31 59 24 36 56 38 47 48 57 24 ...
 $ Gender           : Factor w/ 2 levels "Female","Male": 2 1 1 1 2 1 2 2 2 1 ...
 $ Years.at.Company : int   19 4 10 7 41 3 23 16 44 1 ...
 $ Job.Role         : Factor w/ 5 levels "Education","Finance",...: 1 4 3 1 1 5 1 2 1 3 ...
 $ Monthly.Income   : int  5390 5534 8159 3989 4821 9977 3681 11223 3773 7319 ...
 $ Work.Life.Balance : Factor w/ 4 levels "Excellent","Fair",...: 1 4 3 3 2 2 2 1 3 4 ...

```

```

$ Job.Satisfaction      : Factor w/ 4 levels "High","Low","Medium",...: 3 1 1 1 4 1 1 4 3 1 ...
$ Performance.Rating   : Factor w/ 4 levels "Average","Below Average",...: 1 4 4 3 1 2 3 3 3 1 ...
$ Number.of.Promotions : int  2 3 0 1 0 3 1 2 1 1 ...
$ Overtime             : Factor w/ 2 levels "No","Yes": 1 1 1 1 2 1 2 1 2 2 ...
$ Distance.from.Home   : int  22 21 11 27 71 37 75 5 39 57 ...
$ Education.Level      : Factor w/ 5 levels "Associate Degree",...: 1 4 2 3 3 2 3 4 3 5 ...
$ Marital.Status       : Factor w/ 3 levels "Divorced","Married",...: 2 1 2 3 1 2 1 2 2 3 ...
$ Number.of.Dependents : int  0 3 3 2 0 0 3 4 4 4 ...
$ Job.Level            : Factor w/ 3 levels "Entry","Mid",...: 2 2 2 2 3 2 1 1 1 1 ...
$ Company.Size         : Factor w/ 3 levels "Large","Medium",...: 2 2 2 3 2 2 3 2 2 1 ...
$ Company.Tenure       : int  89 21 74 50 68 47 93 88 75 45 ...
$ Remote.Work          : Factor w/ 2 levels "No","Yes": 1 1 1 2 1 1 1 1 1 1 ...
$ Leadership.Opportunities: Factor w/ 2 levels "No","Yes": 1 1 1 1 1 1 1 1 1 1 ...
$ Innovation.Opportunities: Factor w/ 2 levels "No","Yes": 1 1 1 1 1 2 1 1 1 2 ...
$ Company.Reputation    : Factor w/ 4 levels "Excellent","Fair",...: 1 2 4 3 2 2 3 1 2 3 ...
$ Employee.Recognition  : Factor w/ 4 levels "High","Low","Medium",...: 3 2 2 3 3 1 3 2 3 2 ...
$ Attrition             : Factor w/ 2 levels "Left","Stayed": 2 2 2 2 2 1 1 2 2 1 ...

```

## Data Preprocessing

To prepare the dataset for further analysis, several data preprocessing steps are performed:

### 1. Removing features

```

# first column contains Employee IDs, so not necessary for analysis
data <- data[, !names(data) %in% "Employee.ID"]

# EXPLANATION!!!!!!!!!!!!
data <- data[, !names(data) %in% "Company.Tenure"]

```

### 2. Numeric and categorical value separation

```

numeric_vars <- sapply(data, is.numeric)
categoric_vars <- sapply(data, function(x) is.factor(x) || is.character(x))

```

### 3. Handling missing values

```

# Missing Values --- No null Values
sapply(data, function(x) sum(is.na(x)))

```

Age	Gender	Years.at.Company
0	0	0
Job.Role	Monthly.Income	Work.Life.Balance
0	0	0
Job.Satisfaction	Performance.Rating	Number.of.Promotions
0	0	0
Overtime	Distance.from.Home	Education.Level
0	0	0

Marital.Status	Number.of.Dependents	Job.Level
0	0	0
Company.Size	Remote.Work	Leadership.Opportunities
0	0	0
Innovation.Opportunities	Company.Reputation	Employee.Recognition
0	0	0
Attrition		
0		

## Categorical Features

```
# Categorical feature names
categoric_var_names <- names(data)[categoric_vars]

# Categorical value distributions
for (var in categoric_var_names) {
  cat("\nDistribution of", var, ":\n")
  print(table(data[[var]]))
}
```

Distribution of Gender :

Female	Male
33672	40826

Distribution of Job.Role :

Education	Finance	Healthcare	Media	Technology
15658	10448	17074	11996	19322

Distribution of Work.Life.Balance :

Excellent	Fair	Good	Poor
13432	22529	28158	10379

Distribution of Job.Satisfaction :

High	Low	Medium	Very High
37245	7457	14717	15079

Distribution of Performance.Rating :

Average	Below Average	High	Low
44719	11139	14910	3730

Distribution of Overtime :

No Yes  
50157 24341

Distribution of Education.Level :

Associate Degree	Bachelor's Degree	High School	Master's Degree
18649	22331	14680	15021
PhD			
3817			

Distribution of Marital.Status :

Divorced	Married	Single
11078	37419	26001

Distribution of Job.Level :

Entry	Mid	Senior
29780	29678	15040

Distribution of Company.Size :

Large	Medium	Small
14912	37231	22355

Distribution of Remote.Work :

No Yes  
60300 14198

Distribution of Leadership.Opportunities :

No Yes  
70845 3653

Distribution of Innovation.Opportunities :

No Yes  
62394 12104

Distribution of Company.Reputation :

Excellent	Fair	Good	Poor
7414	14786	37182	15116

Distribution of Employee.Recognition :

High	Low	Medium	Very High
------	-----	--------	-----------



18550      29620      22657      3671











Distribution of Attrition :

Left Stayed

35370 39128

```
# Categorical value distribution -- barplot
par(mfrow = c(2, 3))
for (cat_var in categoric_var_names) {
  barplot(table(data[[cat_var]]), main = paste(cat_var, "Distribution"),
    xlab = cat_var, col = "firebrick4")
}
```




**Education.Level Distribution****Marital.Status Distribution****Job.Level Distribution****Company.Size Distribution****Remote.Work Distribution****Leadership.Opportunities Distribution****Innovation.Opportunities Distribution****Company.Reputation Distribution****Employee.Recognition Distribution****Attrition Distribution**

## Target Values

```
# Target value distribution
par(mfrow = c(1, 2))
barplot(table(data$Attrition), main = "Attrition Count", xlab = "Attrition",
        ylab = "Count", col = c("firebrick4", "rosybrown2"))

# Target value distribution -- Pie chart
attrition_table <- table(data$Attrition)
attrition_df <- as.data.frame(attrition_table)
colnames(attrition_df) <- c("Attrition", "Count")
attrition_df$Percentage <- round(100 * attrition_df$Count/sum(attrition_df$Count),
1)

pie(attrition_df$Count, labels = paste(attrition_df$Attrition, " - ",
↵ attrition_df$Percentage,
"%"), col = c("firebrick4", "rosybrown2"), main = "Attrition Distribution",
cex = 1.2, radius = 0.8)
```



## Numeric Features

```
numeric_var_names <- names(data)[numeric_vars]

# Copy Data
data_detailed <- data


# Define bins for numeric values
data_detailed$Age_Cat <- cut(data_detailed$Age, breaks = c(18, 25, 35, 45,
  55, 60), labels = c("18-25", "25-35", "35-45", "45-55", "55-60"), right = FALSE)
data_detailed$MonthlyIncome_Cat <- cut(data_detailed$Monthly.Income, breaks = c(0,
  3000, 6000, 9000, 12000, 15000, 18000), labels = c("0-3000", "3000-6000",
  "6000-9000", "9000-12000", "12000-15000", "15000+"), right = FALSE)
data_detailed$DistanceFromHome_Cat <- cut(data_detailed$Distance.from.Home,
  breaks = c(0, 20, 40, 60, 80, 100), labels = c("0-20", "20-40", "40-60",
  "60-80", "80-100"), right = FALSE)
data_detailed$YearsAtCompany_Cat <- cut(data_detailed$Years.at.Company, breaks = c(0,
  10, 20, 30, 40, 50, 60), labels = c("0-10", "10-20", "20-30", "30-40",
  "40-50", "50-60"), right = FALSE)

exclude_columns <- c("Age", "Monthly.Income", "Distance.from.Home", "Years.at.Company",
  "Attrition")
data_corr <- data_detailed[, !names(data_detailed) %in% exclude_columns]
data_corr <- data_corr[, c(setdiff(names(data_corr), "Attrition"))]


# Plotting with target feature after transform numeric features into
# categorical
par(mfrow = c(3, 4))
for (col in setdiff(names(data_detailed), exclude_columns)) {
  table_left <- table(data_detailed[data_detailed$Attrition == "Left",
    col])
  table_stayed <- table(data_detailed[data_detailed$Attrition == "Stayed",
    col])

  barplot(rbind(table_left, table_stayed), beside = TRUE, main = paste(col,
    "Distribution by Attrition"), xlab = col, col = c("firebrick4", "rosybrown2"))
}
```

### Gender Distribution by Attb.Role Distribution by Atife.Balance Distribution by Job.Satisfaction Distribution by




### Performance.Rating Distribution by Number.of.Promotions Distribution by Overtime Distribution by Education.Level Distribution by




### Marital.Status Distribution by Number.of.Dependents Distribution by Job.Level Distribution by Company.Size Distribution by




### Remote.Work Distribution by Leadership.Opportunities Distribution by Innovation.Opportunities Distribution by Company.Reputation Distribution by



### Employee.Recognition Distribution by Age\_Cat Distribution by MonthlyIncome\_Cat Distribution by DistanceFromHome\_Cat Distribution by



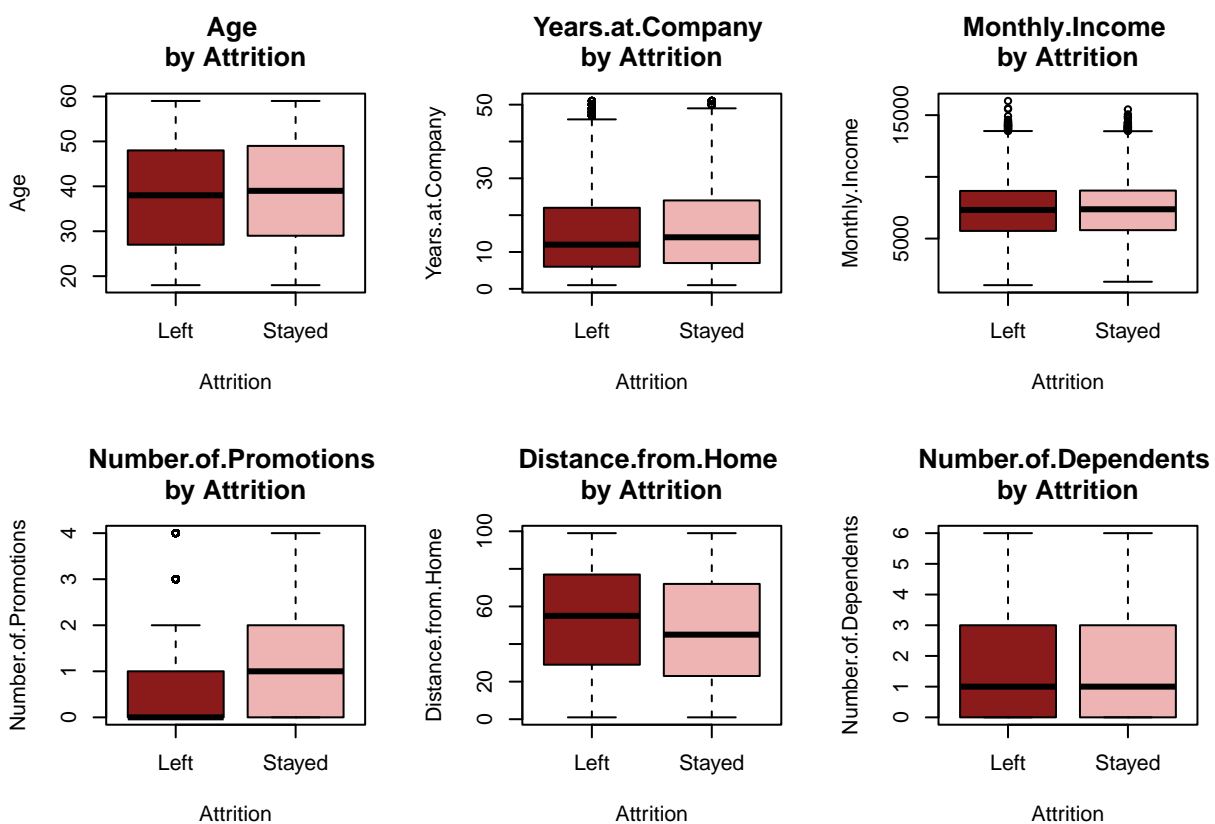
### YearsAtCompany\_Cat Distribution



## Outliers

```
# Target Visualization with Numeric features Outlier Check -- boxplot
cat_var <- "Attrition"

plots_per_page <- 6
num_plots <- length(numeric_var_names)
num_pages <- ceiling(num_plots/plots_per_page)
plot_index <- 1
for (page in 1:num_pages) {
  par(mfrow = c(2, 3))
  for (i in 1:plots_per_page) {
    if (plot_index > num_plots)
      break
    num_var <- numeric_var_names[plot_index]
    boxplot(data[[num_var]] ~ data[[cat_var]], main = paste(num_var,
      "\nby", cat_var), xlab = cat_var, ylab = num_var, col = c("firebrick4",
      "rosybrown2"))
    plot_index <- plot_index + 1
  }
}
```



```
# Function to identify outliers using IQR
identify_outliers <- function(x) {
  Q1 <- quantile(x, 0.25, na.rm = TRUE)
  Q3 <- quantile(x, 0.75, na.rm = TRUE)
  IQR <- Q3 - Q1
  outliers <- x[x < (Q1 - 1.5 * IQR) | x > (Q3 + 1.5 * IQR)]
  return(outliers)
}

# Identify and show outliers for each numeric variable
outliers_list <- list()
for (var in numeric_var_names) {
  outliers <- identify_outliers(data[[var]])
  outliers_list[[var]] <- outliers
  cat("\nOutlier values in", var, "are between:\n")
  print(sort(unique(outliers)))
}
```

Outlier values in Age are between:  
integer(0)

Outlier values in Years.at.Company are between:  
[1] 48 49 50 51

Outlier values in Monthly.Income are between:  
[1] 13713 13715 13722 13728 13739 13747 13764 13768 13793 13800 13830 13833  
[13] 13840 13859 13875 13876 13877 13893 13896 13959 13961 13962 13988 14002  
[25] 14014 14016 14021 14028 14066 14067 14076 14103 14110 14127 14137 14147  
[37] 14163 14176 14181 14185 14210 14226 14235 14271 14276 14286 14292 14396  
[49] 14398 14405 14406 14412 14421 14433 14547 14622 14707 14839 14885 14924  
[61] 15063 15464 15495 15552 16149

Outlier values in Number.of.Promotions are between:  
integer(0)

Outlier values in Distance.from.Home are between:  
integer(0)

Outlier values in Number.of.Dependents are between:  
integer(0)

```
# Plot histograms and highlight outliers

plot_index <- 1
for (page in 1:num_pages) {
  par(mfrow = c(2, 3))
  for (i in 1:plots_per_page) {
```

```

if (plot_index > num_plots)
  break
num_var <- numeric_var_names[plot_index]

hist(data[[num_var]], main = paste(num_var, "Distribution"), xlab = num_var,
      col = "firebrick4", breaks = 30)
outliers <- outliers_list[[num_var]]
if (length(outliers) > 0) {
  points(outliers, rep(0, length(outliers)), col = "rosybrown2",
        pch = 16)
}

plot_index <- plot_index + 1
}
}

```



As a result of the analysis, the following observations were made regarding the characteristics of the data:

- 
- 
- 
- 
- 
-



.

## Features vs. Target

### Categorical Features vs. Target

### Numerical Features vs. Target

## Correlation Matrix

*# Cor. and Cov.*

```
cov_matrix <- cov(data[, numeric_var_names])
cor_matrix <- cor(data[, numeric_var_names])
```

```
print("Covariance Matrix:")
```

```
[1] "Covariance Matrix:"
```

```
print(cov_matrix)
```

	Age	Years.at.Company	Monthly.Income
Age	146.009914270	72.87199387	-45.51579
Years.at.Company	72.871993868	125.97242911	-144.24846
Monthly.Income	-45.515785898	-144.24846489	4633293.12554
Number.of.Promotions	0.008083759	0.01048464	12.14436
Distance.from.Home	-1.579927372	-1.54734492	-117.21026
Number.of.Dependents	0.069262885	0.07649657	5.04010

	Number.of.Promotions	Distance.from.Home
Age	0.008083759	-1.57992737
Years.at.Company	0.010484640	-1.54734492
Monthly.Income	12.144360519	-117.21026410
Number.of.Promotions	0.990599761	-0.19392391
Distance.from.Home	-0.193923912	813.02599733
Number.of.Dependents	-0.002255666	-0.04226003

	Number.of.Dependents
Age	0.069262885
Years.at.Company	0.076496571
Monthly.Income	5.040099975
Number.of.Promotions	-0.002255666
Distance.from.Home	-0.042260030
Number.of.Dependents	2.413775012

```
print("Correlation Matrix:")
```

```
[1] "Correlation Matrix:"
```

```
print(cor_matrix)
```

	Age	Years.at.Company	Monthly.Income
Age	1.0000000000	0.537318418	-0.001749951
Years.at.Company	0.5373184182	1.000000000	-0.005970745
Monthly.Income	-0.0017499514	-0.005970745	1.000000000
Number.of.Promotions	0.0006721606	0.000938570	0.005668663
Distance.from.Home	-0.0045855743	-0.004835008	-0.001909715
Number.of.Dependents	0.0036894448	0.004386881	0.001507113

	Number.of.Promotions	Distance.from.Home
Age	0.0006721606	-0.0045855743
Years.at.Company	0.0009385700	-0.0048350077
Monthly.Income	0.0056686632	-0.0019097149
Number.of.Promotions	1.0000000000	-0.0068332929
Distance.from.Home	-0.0068332929	1.0000000000
Number.of.Dependents	-0.0014587377	-0.0009539579

	Number.of.Dependents
Age	0.0036894448
Years.at.Company	0.0043868807
Monthly.Income	0.0015071132
Number.of.Promotions	-0.0014587377
Distance.from.Home	-0.0009539579
Number.of.Dependents	1.0000000000

```
corrplot(cor_matrix, method = "number", type = "upper", tl.col = "black",
         tl.srt = 45)
```



### Partial Correlation Matrices

## Data Preparation

After completing the data analysis steps, it is necessary to prepare the data for model development.

### Handling Categorical Features

In order to use the categorical features in the model, we need to convert categorical features to numeric (ordinal or nominal) representations.

```
# Ordinal mappings:
balance.map <- c(Poor = 1, Fair = 2, Good = 3, Excellent = 4)
data$Work.Life.Balance <- balance.map[as.numeric(data$Work.Life.Balance)]

satisfaction.map <- c(Low = 1, Medium = 2, High = 3, `Very High` = 4)
data$Job.Satisfaction <- satisfaction.map[as.numeric(data$Job.Satisfaction)]

performance.map <- c(Low = 1, `Below Average` = 2, Average = 3, High = 4)
data$Performance.Rating <- performance.map[as.numeric(data$Performance.Rating)]
```

```

level.map <- c(Entry = 1, Mid = 2, Senior = 3)
data$Job.Level <- level.map[as.numeric(data$Job.Level)]

# education.map <- c('High School'= 1, 'Associate Degree'= 2,
# 'Bachelor's Degree'= 3, 'Master's Degree'= 4, 'PhD'= 5)
# data$Education.Level <-
# education.map[as.numeric(data$Education.Level)]

reputation.map <- c(Poor = 1, Fair = 2, Good = 3, Excellent = 4)
data$Company.Reputation <- reputation.map[as.numeric(data$Company.Reputation)]

recognition.map <- c(Low = 1, Medium = 2, High = 3, `Very High` = 4)
data$Employee.Recognition <- recognition.map[as.numeric(data$Employee.Recognition)]

size.map <- c(Small = 1, Medium = 2, Large = 3)
data$Company.Size <- size.map[as.numeric(data$Company.Size)]

# Nominal mappings: Create dummy variables for nominal data
data_numeric <- model.matrix(~., data = data)

# Convert the resulting matrix back to a data frame
data_numeric <- as.data.frame(data_numeric)[, -1] # -1 to remove the intercept column

```

## Normalization

```

normalize <- function(x) {
  return((x - min(x))/(max(x) - min(x)))
}

data_normalized <- as.data.frame(lapply(data_numeric, normalize))

```

## Train-Test-Split

Before splitting the data into training and test, first features and target should be defined.

```

# Splitting data into features and target:
X <- data_normalized[, !(colnames(data_normalized) %in% c("Employee.ID",
  "AttritionStayed"))]

y <- data_normalized$AttritionStayed

```

Now, we can split the dataset for modelling.

```
set.seed(123)

trainIndex <- sample(1:nrow(X), 0.8 * nrow(X))

# 80% of data is used for training
X.train <- X[trainIndex, ]
y.train <- y[trainIndex]

# 20% of data is used for testing
X.test <- X[-trainIndex, ]
y.test <- y[-trainIndex]
```

Before moving to modelling step, it is beneficial to check the dimensions and balance of the datasets.

- Number of samples in train data: 59598
  - Proportion of stayed employees for train data:

```
kable(prop.table(table(y.train)))
```

y.train	Freq
0	0.4733213
1	0.5266787

- Number of samples in test data: 14900
  - Proportion of stayed employees for test data:

```
prop.table(table(y.test))
```

```
y.test
      0      1
0.480604 0.519396
```

We can observe that the train and test datasets are balanced within themselves. Also the train data is representative of test data.

## Predictive Classification Models

Predictive classification models are a type of machine learning algorithm used to predict the category or class label of new, unseen instances based on historical data. These models are trained using a labelled dataset where the input features (independent variables) are associated with known class labels (dependent variable). The goal of the model is to learn the relationship between the features and the class labels so that it can accurately classify new data points into one of the predefined categories.

In this project we aim to find the risk of an employee leaving the company (class 0) and the factors affecting employee retention. So we will develop several classification models and examine their performances.

## Logistic Regression

The logistic regression model estimates the odds of the dependent variable occurring and applies the logit (log-odds) transformation to express this relationship.

### Basic Logistic Classifier

```
# First of all we check the model statistics with all the features
glm.FULL <- glm(y.train ~ ., data = X.train, family = binomial)
summary(glm.FULL)
```

Call:

```
glm(formula = y.train ~ ., family = binomial, data = X.train)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-0.553730	0.064437	-8.593	< 2e-16	***
Age	0.229757	0.039315	5.844	5.10e-09	***
GenderMale	0.543694	0.019788	27.476	< 2e-16	***
Years.at.Company	0.654574	0.051922	12.607	< 2e-16	***
Job.RoleFinance	0.109910	0.046309	2.373	0.017625	*
Job.RoleHealthcare	0.061875	0.040506	1.528	0.126628	
Job.RoleMedia	0.106027	0.034483	3.075	0.002107	**
Job.RoleTechnology	0.098439	0.046401	2.121	0.033880	*
Monthly.Income	0.018438	0.117791	0.157	0.875614	
Work.Life.Balance	-0.565409	0.031321	-18.052	< 2e-16	***
Job.Satisfaction	-0.365624	0.024033	-15.213	< 2e-16	***
Performance.Rating	-0.289214	0.030814	-9.386	< 2e-16	***
Number.of.Promotions	0.928255	0.039924	23.250	< 2e-16	***
OvertimeYes	-0.336407	0.020902	-16.095	< 2e-16	***
Distance.from.Home	-0.886487	0.033969	-26.097	< 2e-16	***
Education.LevelBachelor.s.Degree	-0.035818	0.026276	-1.363	0.172835	
Education.LevelHigh.School	0.001582	0.029370	0.054	0.957050	
Education.LevelMaster.s.Degree	-0.006862	0.029087	-0.236	0.813509	
Education.LevelPhD	1.506973	0.053149	28.354	< 2e-16	***
Marital.StatusMarried	0.257001	0.028268	9.092	< 2e-16	***
Marital.StatusSingle	-1.409863	0.030515	-46.203	< 2e-16	***
Number.of.Dependents	0.842765	0.038204	22.060	< 2e-16	***
Job.Level	2.292477	0.028636	80.056	< 2e-16	***
Company.Size	-0.193325	0.027989	-6.907	4.94e-12	***
Remote.WorkYes	1.635385	0.027848	58.726	< 2e-16	***
Leadership.OpportunitiesYes	0.162824	0.045057	3.614	0.000302	***
Innovation.OpportunitiesYes	0.127904	0.026574	4.813	1.49e-06	***
Company.Reputation	-0.342783	0.033757	-10.154	< 2e-16	***
Employee.Recognition	-0.003910	0.034468	-0.113	0.909677	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 82451 on 59597 degrees of freedom  
 Residual deviance: 62455 on 59569 degrees of freedom  
 AIC: 62513

Number of Fisher Scoring iterations: 4

The above model statistics indicate that p-value of Employee Recognition is above 0.5 indicating that this feature is insignificant to the results. Additionally, some Job.Roles and Monthly. Income also have high p-values indicating that their effect on Attrition is less significant compared to other features. However, for now we would like to keep all the features in the model and apply feature selection later.

In order to understand how well the model fits the data we can make use of  $R^2$  statistics.  $R^2$  provides an indication of how well the independent variables in the model explain the variability of the dependent variable. A higher  $R^2$  value indicates a better fit of the model to the data. The formula for  $R^2$  is:

$$R^2 = 1 - \frac{RSS}{ESS}$$

Where:

- $RSS$  is the sum of squares of the residuals (the differences between observed and predicted values), i.e. the deviance of the fitted model
- $ESS$  is the total sum of squares due to regression (the differences between the observed values and the mean of the observed values)

```
R2 <- 1 - (summary(glm.FULL)$deviance/summary(glm.FULL)$null.deviance)
R2
```

```
[1] 0.2352223
```

With the full model the value of  $R^2$  0.2352223 indicates that approximately 23.5222337% of the variance in the target can be explained by the features in the model. Since 23.5222337% is relatively low, it suggests that the model is not capturing much of the underlying pattern in the data.

Multicollinearity can be a reason for a low  $R^2$  value, as it can make it difficult to determine the individual effect of each predictor on the target. Calculating the Variance Inflation Factor (VIF) can help to check for multicollinearity among the features.

```
vif.FULL <- data.frame(features = names(vif(glm.FULL)), VIF = vif(glm.FULL),
  row.names = NULL)
vif.FULL[order(-vif.FULL$VIF), ]
```

	features	VIF
7	Job.RoleTechnology	4.303457
5	Job.RoleHealthcare	3.012165
8	Monthly.Income	2.998615
4	Job.RoleFinance	2.698510

17	Marital.StatusSingle	2.158750
16	Marital.StatusMarried	2.081797
6	Job.RoleMedia	1.678546
3	Years.at.Company	1.404475
1	Age	1.401725
19	Job.Level	1.093697
21	Remote.WorkYes	1.058167
2	GenderMale	1.013838
14	Distance.from.Home	1.009712
12	Number.of.Promotions	1.009666
18	Number.of.Dependents	1.008320
9	Work.Life.Balance	1.006144
13	OvertimeYes	1.005366
10	Job.Satisfaction	1.004918
15	Education.Level	1.004629
24	Company.Reputation	1.002553
11	Performance.Rating	1.001966
20	Company.Size	1.001456
22	Leadership.OpportunitiesYes	1.000757
23	Innovation.OpportunitiesYes	1.000466
25	Employee.Recognition	1.000405

A VIF value of 1 indicates no correlation, values between 1 and 5 indicate moderate correlation and values above 5 suggest significant multicollinearity, which can lead to unreliable coefficient estimates. Most variables have VIF values close to 1, indicating very low or no multicollinearity. A few variables have VIF values between 1 and 5, suggesting moderate multicollinearity, which may not pose serious issues but should be monitored. These variables include Job.RoleTechnology, Job.RoleHealthcare, Monthly.Income, Job.RoleFinance, Marital.StatusSingle and Marital.StatusMarried. These are mostly dummy features of nominal variables and dummy variables are often correlated because they represent categories of the same nominal variable.

### Logistic Regression with Backward Stepwise Search

Backward variable selection is a greedy search algorithm used to develop a predictive model by iteratively removing the least significant features. The goal is to find the best subset of features that contribute to the model while eliminating those that do not improve its performance.

We can use the `step()` function to perform backward stepwise search. As a regularization criteria we can either use BIC or AIC. But BIC statistic generally places a heavier penalty on models with many variables, and hence results in the selection of smaller models than AIC.

```
# Backward Stepwise Search with AIC statistics
glm.BACKWARD <- step(glm.FULL, direction = "backward", k = log(nrow(X.train)),
  trace = 0, steps = 100)
summary(glm.BACKWARD)
```

Call:

```
glm(formula = y.train ~ Age + GenderMale + Years.at.Company +
```



```
Work.Life.Balance + Job.Satisfaction + Performance.Rating +
Number.of.Promotions + OvertimeYes + Distance.from.Home +
Education.LevelPhD + Marital.StatusMarried + Marital.StatusSingle +
Number.of.Dependents + Job.Level + Company.Size + Remote.WorkYes +
Leadership.OpportunitiesYes + Innovation.OpportunitiesYes +
Company.Reputation, family = binomial, data = X.train)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-0.48883	0.05177	-9.442	< 2e-16	***
Age	0.22990	0.03930	5.849	4.94e-09	***
GenderMale	0.54243	0.01978	27.423	< 2e-16	***
Years.at.Company	0.65402	0.05191	12.599	< 2e-16	***
Work.Life.Balance	-0.56426	0.03131	-18.022	< 2e-16	***
Job.Satisfaction	-0.36546	0.02403	-15.210	< 2e-16	***
Performance.Rating	-0.28983	0.03080	-9.409	< 2e-16	***
Number.of.Promotions	0.92821	0.03991	23.256	< 2e-16	***
OvertimeYes	-0.33570	0.02089	-16.068	< 2e-16	***
Distance.from.Home	-0.88551	0.03396	-26.076	< 2e-16	***
Education.LevelPhD	1.51895	0.05049	30.086	< 2e-16	***
Marital.StatusMarried	0.25788	0.02826	9.126	< 2e-16	***
Marital.StatusSingle	-1.40864	0.03050	-46.183	< 2e-16	***
Number.of.Dependents	0.84234	0.03820	22.052	< 2e-16	***
Job.Level	2.29026	0.02862	80.023	< 2e-16	***
Company.Size	-0.19248	0.02798	-6.879	6.02e-12	***
Remote.WorkYes	1.63481	0.02784	58.729	< 2e-16	***
Leadership.OpportunitiesYes	0.16201	0.04504	3.597	0.000322	***
Innovation.OpportunitiesYes	0.12815	0.02657	4.823	1.41e-06	***
Company.Reputation	-0.34247	0.03374	-10.149	< 2e-16	***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 82451 on 59597 degrees of freedom  
Residual deviance: 62476 on 59578 degrees of freedom  
AIC: 62516

Number of Fisher Scoring iterations: 4

```
vif.BACKWARD <- data.frame(features = names(vif(glm.BACKWARD)), VIF = vif(glm.BACKWARD),
  row.names = NULL)
vif.BACKWARD[order(-vif.BACKWARD$VIF), ]
```

	features	VIF
12	Marital.StatusSingle	2.163022
11	Marital.StatusMarried	2.085498
3	Years.at.Company	1.405299
1	Age	1.402928

```

14          Job.Level 1.097342
16      Remote.WorkYes 1.060083
10      Education.LevelPhD 1.015134
2          GenderMale 1.012783
9      Distance.from.Home 1.011547
7      Number.of.Promotions 1.009742
13      Number.of.Dependents 1.009113
4          Work.Life.Balance 1.006760
5          Job.Satisfaction 1.005193
8          OvertimeYes 1.004790
19      Company.Reputation 1.002462
6      Performance.Rating 1.001593
15      Company.Size 1.001324
17 Leadership.OpportunitiesYes 1.000616
18 Innovation.OpportunitiesYes 1.000482

```

It looks like backward stepwise search dropped “Job.RoleFinance”, “Job.RoleHealthcare”, “Job.RoleTechnology” and “Employee.Recognition” from the model. These features were also the ones with highest p-values so it seems logical. But oddly the model kept the “Job.RoleMedia” feature so this tells us that having a job role in Media may be more significant to employee Attrition than other job roles for this dataset.

With backward feature selection we were able to decrease the VIF scores, p-values and AIC score. Although  $R^2$  slightly decreased to 0.2422558, this is expectable since the number of features in the model decreased. But most importantly we decreased the variance inflation factor of the model so the model is more stable.

### Logistic Regression with Shrinkage Method

Shrinkage methods are techniques used in regression analysis to prevent overfitting by introducing a penalty for large coefficients. These methods “shrink” the coefficients towards zero, effectively reducing their variance and, in turn, enhancing the model’s generalizability. Two most used shrinkage methods are Ridge Regression and Lasso Regression.

- Ridge Regression adds a penalty to the regression model that shrinks all the coefficients, and is useful when we want to keep all features in the model.
- Lasso Regression introduces a penalty that can shrink some coefficients to zero. This method is useful for feature selection, yielding sparse models by retaining only the most important features.
- Elastic Net is a hybrid regularization technique that includes both the Ridge and Lasso penalties, allowing for variable selection (like Lasso) and handling multicollinearity (like Ridge). It aims to incorporate the strengths of both methods, providing a more flexible approach.


Choosing the value of  $\lambda$  (the tuning parameter) is crucial because it determines the strength of the penalty applied to the coefficients. We can choose the best value of  $\lambda$  using k-fold cross-validation method.

```

# Elastic Net Regression with alpha=0.05 (more weight on Ridge Reg.)
set.seed(123)

glm.ELNET <- cv.glmnet(as.matrix(X.train), y.train, alpha = 0.05, family = "binomial",
  type.measure = "class")
best.lamda <- glm.ELNET$lambda.min
plot(glm.ELNET)

```



### Comparison of Logistic Classifiers


For Logistic Regression we defined 3 Logistic classifiers. In order to identify the best model we can compare their performance on the test sets to see how well they captured the underlying pattern of the data and their ability to generalize to new data.

```
# 1. Basic Logistic Classifier
glm.FULL.predict <- predict(glm.FULL, newdata = X.test, type = "response")

# 2. Feature Selection with Backward Stepwise Search
glm.BACKWARD.predict <- predict(glm.BACKWARD, newdata = X.test, type = "response")

# 3. Elastic Net Shrinkage Method
glm.ELNET.predict <- predict(glm.ELNET, newx = as.matrix(X.test), type = "response",
  s = best.lamda)
```

We can use ROC curve and AUC values to compare the models. The ROC curve is a tool for assessing the performance of binary classification models, plotting true positive rate against false positive rate at various thresholds. The Area Under the Curve (AUC) provides a measure of the model's ability to predict the target values, with higher values indicating better performance.



Although the AUC values are same, to better analyse the difference between the model performances we can calculate and compare the evaluation metrics.

```
# Function to evaluate prediction models at different thresholds
evaluate.model.performance <- function(predictions, y.test, thresholds, title) {
  output.list <- list()

  # Looping through each threshold to generate predictions and store
  # in output.list
  for (threshold in thresholds) {
    output <- ifelse(predictions > threshold, 1, 0)
    output.list[[as.character(threshold)]] <- output
  }

  # Initialize a data frame to store the evaluation metrics for each
  # threshold
  results <- data.frame(Threshold = numeric(length(thresholds)), Accuracy =
    numeric(length(thresholds)),
    F1.Score = numeric(length(thresholds)), Precision = numeric(length(thresholds)),
    Recall = numeric(length(thresholds)))

  # Compute evaluation metrics for each threshold
  for (i in 1:length(thresholds)) {
    threshold <- thresholds[i]
```

```

predict.output <- output.list[[as.character(threshold)]]

# Store results in the results dataframe
results[i, "Threshold"] = threshold
results[i, "Accuracy"] = mean(predict.output == y.test)
results[i, "F1.Score"] = (2 * sum((predict.output == 1 & y.test ==
  1))/(2 * sum((predict.output == 1 & y.test == 1)) + sum(predict.output ==
  1 & y.test == 0) + sum(predict.output == 0 & y.test == 1)))
results[i, "Precision"] = sum(predict.output == 1 & y.test ==
  1)/sum(predict.output ==
  1)
results[i, "Recall"] = sum(predict.output == 1 & y.test == 1)/sum(y.test ==
  1)
}

# Rounding results
results[, c("Accuracy", "F1.Score", "Precision", "Recall")] <- round(results[,
  c("Accuracy", "F1.Score", "Precision", "Recall")], 4)

# Return the results table
kable(results, align = "c", caption = c("Performance Metrics", title))
}

```

```
thresholds <- c(0.2, 0.3, 0.4, 0.5, 0.6)
```

```
# 1. Basic Logistic Classifier
```

```
evaluate.model.performance(glm.FULL.predict, y.test, thresholds, "glm.FULL")
```

**Table 2: Performance Metrics Table: glm.FULL**

Threshold	Accuracy	F1.Score	Precision	Recall
0.2	0.5168	0.6477	0.5213	0.8551
0.3	0.5135	0.6149	0.5221	0.7479
0.4	0.5085	0.5749	0.5219	0.6399
0.5	0.5033	0.5256	0.5215	0.5298
0.6	0.5009	0.4682	0.5243	0.4229

```
# 2. Feature Selection with Backward Stepwise Search
```

```
evaluate.model.performance(glm.BACKWARD.predict, y.test, thresholds, "glm.BACKWARD")
```

**Table 3: Performance Metrics Table: glm.BACKWARD**

Threshold	Accuracy	F1.Score	Precision	Recall
0.2	0.5169	0.6477	0.5213	0.8549
0.3	0.5134	0.6151	0.5220	0.7487
0.4	0.5091	0.5753	0.5224	0.6401
0.5	0.5036	0.5259	0.5218	0.5300
0.6	0.5005	0.4677	0.5238	0.4225

# 3. Elastic Net Shrinkage Method

```
evaluate.model.performance(glm.ELNET.predict, y.test, thresholds, "glm.ELNET")
```

**Table 4: Performance Metrics Table: glm.ELNET**

Threshold	Accuracy	F1.Score	Precision	Recall
0.2	0.5185	0.6549	0.5217	0.8794
0.3	0.5134	0.6222	0.5213	0.7714
0.4	0.5109	0.5822	0.5232	0.6562
0.5	0.5023	0.5253	0.5205	0.5303
0.6	0.5003	0.4612	0.5242	0.4118

All of the models have the highest F1 score at threshold level 0.4. And F1 score with 0.4 threshold for glm.ELNET is higher than other models on the test set. These results indicate that the model with Elastic Net regularization is better in generalization. Therefore we can select the glm.ELNET.predict as the best prediction for logistic regression model.

## Another Classification Model

## Model Results

### Performance Metrics and Confusion Matrix